W3

July 19, 2021

1 Part 1: Visualization Technique

A. Narrative I found the interesting dataset in Kaggle concerning alcohol consumption in the world for adults 15 years and older. I was immediately curious as I pursue a healthy lifestyle. This dataset covers many different regions and countries, and contains many variables that could produce interesting findings.

The Dataset: Alcohol assumption in 2008 downloaded from Kaggle. The original data source came from Gap Minder, a non-profit research entity. The full descriptions of the curated dataset is discussed in detail here: http://makemeanalyst.com/download-and-learn-about-gapminder-dataset/

After having worked on this assignment for a while, I realized that my original dataset is rather limited. It doesn't have interesting layers of categorical data to enhance my visualizations. Hence, I browsed around the web to obtain two additional datasets. These came from the World Health Organization that contains country,region, and income information for the countries of interest. Please note that my demo will attempt to call out this realization and documented the steps chronologically. https://www.who.int/news-room/fact-sheets/detail/alcohol

B. Visualization Techniques:

I will demonstrate some basic charts that were covered in the course, but in this new toolkit (Plotly). Specifically, the charts that I will show are: - Histogram - Bar charts with color coding - Bubble map - Scatter plots (for fun)

C. Discussion:

With this mixed bag of the intended charts above, I will sequentially describe how they work and when they should not be used 1. Histogram - Measures the frequency of a continuous numerical vareiable in the dataset. The bins could be specified - This is not appropriate for categorical data 2. Bar charts with color coding - This is unarguably the most popular visual for anyone. Typtically, one axis is a category variable and the other one shows the numercial data. The data doesn't have to the continuous. I encoded an "extra" feature that shows certain types of grouping as a color callout (see the demo later) 3. Bubble map - I used to work in finance, and bubble chart is used widely to measure sales and their impact across different categories/conditions. And since I have a global dataset, why not combine a bubble and map to create a bubble map? This visual not only effectively describes trends in your data but also offers a "sizing" component (i.e., the bigger the size (coded with a numberical var), the bigger the impact of that data point is with respect to that category. (e.g., Sales rep sales by region and size = total sales made) - It is often not possible to produce this chart if you don't have geographic data for many data points. 4. Scatter plots (extra just for fun) - I would love to understand the relationships between the variables in the dataset.

This chart is often used to understanding the linear correlation between the x and y variables. - x and y must be numberical for this chart to be used.

2 Part 2: Visualization Library

I chose Plotly, an open-source Python library developed by a Canadian firm founded in 2012. The toolkit supports 40 unique plots. It is a more aesthetically pleasing than matplotlib and offers interactive capablity. This comes in handy for my demo as you will see later that it is quite nice to be able to hover your mouse on the data points and glean some data right on the spots! Plotly also highly integrates with Dash, an open-source framework for building analytical applications, with no Javascript required. I will continue to tinker with this library as I progress throughout my journey as a data scientist.

I checked out basic Python documentation via this link:

https://plotly.com/python/

I installed through Conda command line %conda install plotly

[1]:
```
# I will import the crucial standard libraries and modules
import plotly
import plotly.express as px
import pandas as pd
import numpy as np
```

[2]:
```
print(plotly.__version__)
```

```
4.14.3
```

3 Part 3: Demonstration

[3]:
```
# Loading the data, basic manipulations for cleaning and removing unwanted rows
```

[4]:
```
a_data = pd.read_csv('gapminder_alcohol.csv')
print(a_data.shape)
a_data.set_index('country', inplace=True)
a_data.rename(columns={'alcconsumption':'alcconsumption 2008'}, inplace=True)
a_data.head()
```

```
(213, 6)
```

[4]:

country	alcconsumption 2008	incomeperperson	suicideper100th	\
Afghanistan	0.03	NaN	6.684385	
Albania	7.29	1914.996551	7.699330	
Algeria	0.69	2231.993335	4.848770	
Andorra	10.17	21943.339900	5.362179	
Angola	5.57	1381.004268	14.554677	

country	employrate	urbanrate
Afghanistan	55.700001	24.04
Albania	51.400002	46.72
Algeria	50.500000	65.22
Andorra	NaN	88.92
Angola	75.699997	56.70

[5]:
```
# lets explore a basic histogram of alchohol consumption measured in litres for␣
 ↪adults >= 15 years old to see the
# frequency of the different buckets measured in litres
fig = px.histogram(a_data, x="alcconsumption 2008", title='Histogram of Global␣
 ↪Alchohol Consumption per Adult (in litres)')
fig.show()
```

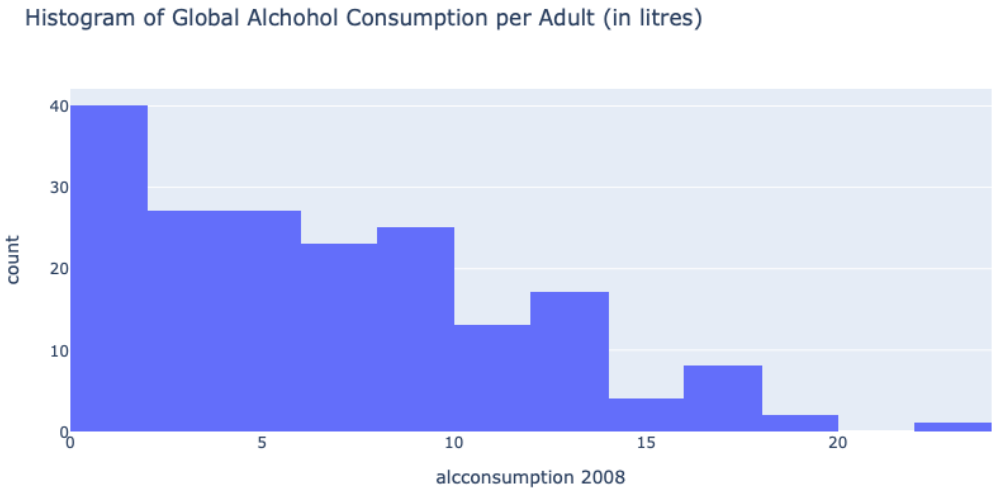

[6]:
```
# The distribution looks right-skewed to me. I'm curious about the basic␣
 ↪summary statistics of the level of alcohol
# consumptions. Let's perform some basic statistical measures.
a_data.describe()
```

[6]:

	alcconsumption 2008	incomeperperson	suicideper100th	employrate	\
count	187.000000	190.000000	191.000000	178.000000	
mean	6.689412	8740.966076	9.640839	58.635955	
std	4.899617	14262.809083	6.300178	10.519454	
min	0.030000	103.775857	0.201449	32.000000	
25%	2.625000	748.245151	4.988449	51.225000	
50%	5.920000	2553.496056	8.262893	58.699999	

```
75%              9.925000      9379.891166        12.328551   64.975000
max             23.010000    105147.437700        35.752872   83.199997
```

```
        urbanrate
count  203.000000
mean    56.769360
std     23.844933
min     10.400000
25%     36.830000
50%     57.940000
75%     74.210000
max    100.000000
```

[7]:
```
# I want to find out the top countries that consume most alcohol in the world
# I will slice a separate df for these top 15 countries
al_df= a_data[['alcconsumption 2008']]
sort_df = al_df.sort_values('alcconsumption 2008', ascending=False)
top20 = sort_df.iloc[:20,:]
top20
```

[7]:
```
                alcconsumption 2008
country
Moldova                       23.01
Korea, Rep.                   19.15
Belarus                       18.85
Ukraine                       17.47
Estonia                       17.24
Czech Rep.                    16.47
Uganda                        16.40
Lithuania                     16.30
Russia                        16.23
Romania                       16.15
Hungary                       16.12
Croatia                       15.00
Slovenia                      14.94
Ireland                       14.92
Poland                        14.43
Portugal                      13.89
Armenia                       13.66
Latvia                        13.45
Azerbaijan                    13.34
Slovak Republic               13.31
```

[8]:
```
fig = px.bar(top20, x = top20.index, y = 'alcconsumption 2008', title = 'Top 20␣
↪countries consumed most alcohol per adult in 2008')
fig.show()
```

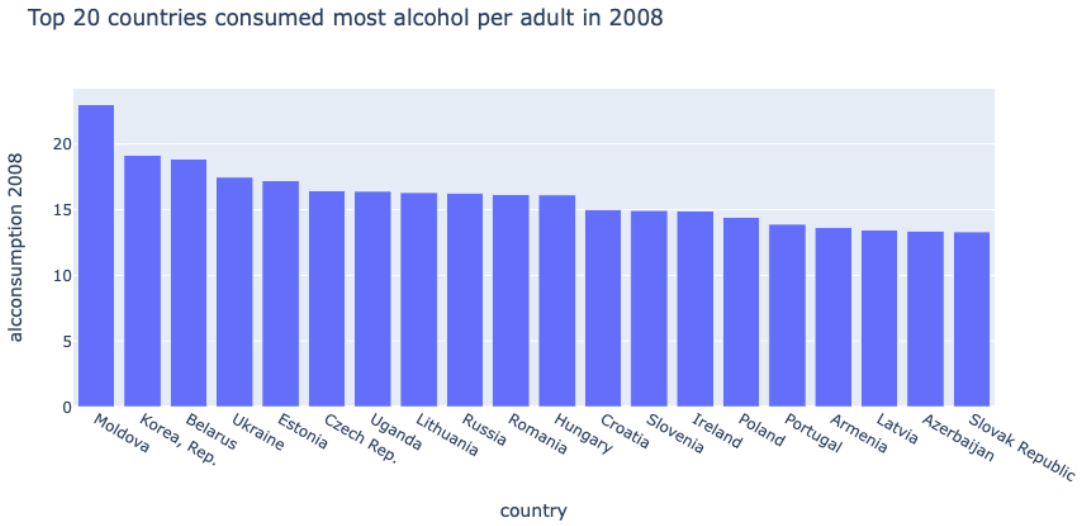

[9]:
```
# now i'm going to bring in two other sources of data to provide an additional dimensions--region and incomegroup so that we could enhance
# this bar chart by turning it into more celebrated bar charts

# 1) country code that will bind these dataframes together (df2)
# 2) region and income information for the countries of interest (df3)
# Finally, i will merge the three dataframes together for a holistic view.
```

[10]:
```
df2 = pd.read_csv('country_code.csv')
df2 = df2.drop(columns=['Indicator Name', 'Indicator Code']) # dropping non-essential columns
print(df2.shape)
df2.head()
```

```
(266, 63)
```

[10]:

	Country Name	Country Code	1960	1961	1962	1963	1964	\
0	Aruba	ABW	NaN	NaN	NaN	NaN	NaN	
1	Africa Eastern and Southern	AFE	NaN	NaN	NaN	NaN	NaN	
2	Afghanistan	AFG	NaN	NaN	NaN	NaN	NaN	
3	Africa Western and Central	AFW	NaN	NaN	NaN	NaN	NaN	
4	Angola	AGO	NaN	NaN	NaN	NaN	NaN	

	1965	1966	1967	…	2011	2012	2013	2014	2015	2016	2017	\
0	NaN	NaN	NaN	…	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	…	NaN	NaN	NaN	NaN	5.200565	NaN	NaN	
2	NaN	NaN	NaN	…	NaN	NaN	NaN	NaN	0.210000	NaN	NaN	
3	NaN	NaN	NaN	…	NaN	NaN	NaN	NaN	6.869468	NaN	NaN	

```
4   NaN   NaN   NaN  …   NaN   NaN   NaN   NaN  7.960000   NaN   NaN

       2018  2019  2020
0       NaN   NaN   NaN
1  5.170911   NaN   NaN
2  0.210000   NaN   NaN
3  6.835266   NaN   NaN
4  6.940000   NaN   NaN

[5 rows x 63 columns]
```

[11]:
```
df3 = pd.read_csv('region_incomegroup.csv')
df3 = df3.drop(columns=['SpecialNotes','TableName', 'Unnamed: 5'])
print(df3.shape)
df3.head()
```

```
(265, 3)
```

[11]:

	Country Code	Region	IncomeGroup
0	ABW	Latin America & Caribbean	High income
1	AFE	NaN	NaN
2	AFG	South Asia	Low income
3	AFW	NaN	NaN
4	AGO	Sub-Saharan Africa	Lower middle income

[12]:
```
# I will merge df2 and df3 together
df_merged = pd.merge(df2, df3,how='outer',left_on='Country Code',␣
↪right_on='Country Code')
print(df_merged.shape)
df_merged.head()
```

```
(266, 65)
```

[12]:

	Country Name	Country Code	1960	1961	1962	1963	1964	\
0	Aruba	ABW	NaN	NaN	NaN	NaN	NaN	
1	Africa Eastern and Southern	AFE	NaN	NaN	NaN	NaN	NaN	
2	Afghanistan	AFG	NaN	NaN	NaN	NaN	NaN	
3	Africa Western and Central	AFW	NaN	NaN	NaN	NaN	NaN	
4	Angola	AGO	NaN	NaN	NaN	NaN	NaN	

	1965	1966	1967	…	2013	2014	2015	2016	2017	2018	2019	\
0	NaN	NaN	NaN	…	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	…	NaN	NaN	5.200565	NaN	NaN	5.170911	NaN	
2	NaN	NaN	NaN	…	NaN	NaN	0.210000	NaN	NaN	0.210000	NaN	
3	NaN	NaN	NaN	…	NaN	NaN	6.869468	NaN	NaN	6.835266	NaN	
4	NaN	NaN	NaN	…	NaN	NaN	7.960000	NaN	NaN	6.940000	NaN	

```
   2020                     Region           IncomeGroup
0   NaN  Latin America & Caribbean          High income
1   NaN                        NaN                  NaN
2   NaN                 South Asia           Low income
3   NaN                        NaN                  NaN
4   NaN        Sub-Saharan Africa  Lower middle income

[5 rows x 65 columns]
```

[13]:
```
# merging all 3 tables together
combined_df = pd.merge(a_data, df_merged,how='inner', left_index=True,␣
↪right_on='Country Name')
combined_df.set_index('Country Name', inplace=True)
print(combined_df.shape)
combined_df.head()
```

```
(182, 69)
```

[13]:

Country Name	alcconsumption 2008	incomeperperson	suicideper100th	…
Afghanistan	0.03	NaN	6.684385	
Albania	7.29	1914.996551	7.699330	
Algeria	0.69	2231.993335	4.848770	
Andorra	10.17	21943.339900	5.362179	
Angola	5.57	1381.004268	14.554677	

Country Name	employrate	urbanrate	Country Code	1960	1961	1962	1963	…	…
Afghanistan	55.700001	24.04	AFG	NaN	NaN	NaN	NaN	…	
Albania	51.400002	46.72	ALB	NaN	NaN	NaN	NaN	…	
Algeria	50.500000	65.22	DZA	NaN	NaN	NaN	NaN	…	
Andorra	NaN	88.92	AND	NaN	NaN	NaN	NaN	…	
Angola	75.699997	56.70	AGO	NaN	NaN	NaN	NaN	…	

Country Name	2013	2014	2015	2016	2017	2018	2019	2020
Afghanistan	NaN	NaN	0.21	NaN	NaN	0.21	NaN	NaN
Albania	NaN	NaN	6.74	NaN	NaN	7.17	NaN	NaN
Algeria	NaN	NaN	0.93	NaN	NaN	0.95	NaN	NaN
Andorra	NaN	NaN	11.01	NaN	NaN	11.02	NaN	NaN
Angola	NaN	NaN	7.96	NaN	NaN	6.94	NaN	NaN

Country Name	Region	IncomeGroup
Afghanistan	South Asia	Low income
Albania	Europe & Central Asia	Upper middle income
Algeria	Middle East & North Africa	Lower middle income

```
Andorra              Europe & Central Asia          High income
Angola                  Sub-Saharan Africa  Lower middle income

[5 rows x 69 columns]
```

[14]:
```
# now i want to revisit my bar chart above and add another level of dimension␣
 ↪to it.
sort_df2 = combined_df.sort_values('alcconsumption 2008', ascending=False)
top20 = sort_df2.iloc[:20,:]

#let's get plotting on this df and add 'region code' and 'income group',␣
 ↪respectively, to encode the colors

fig = px.bar(top20, x = top20.index, y = 'alcconsumption 2008', color='Region',␣
 ↪title = 'Top 20 countries consumed most alcohol in 2008')
fig.show()
fig = px.bar(top20, x = top20.index, y = 'alcconsumption 2008',␣
 ↪color='IncomeGroup', title = 'Top 20 countries consumed most alcohol in␣
 ↪2008')
fig.show()
```

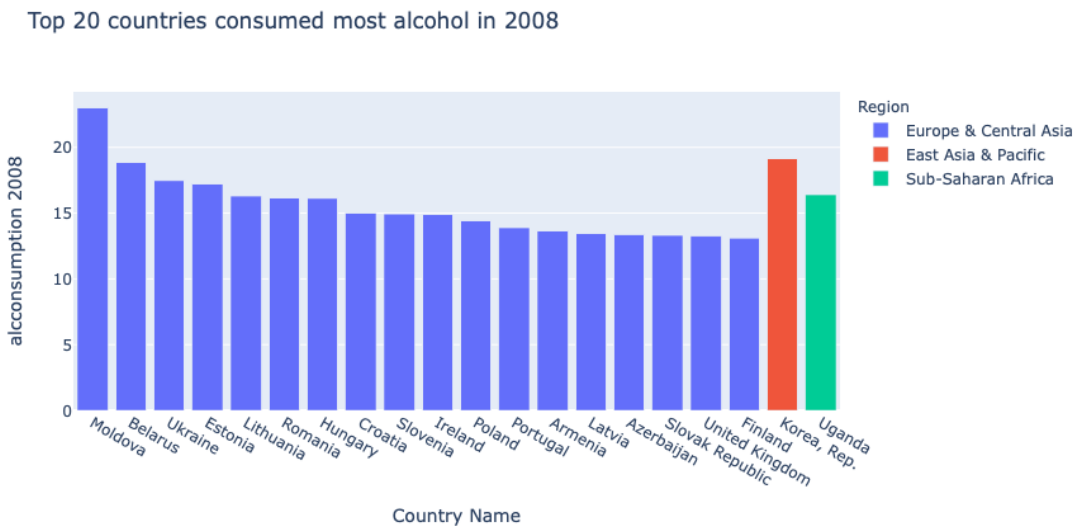

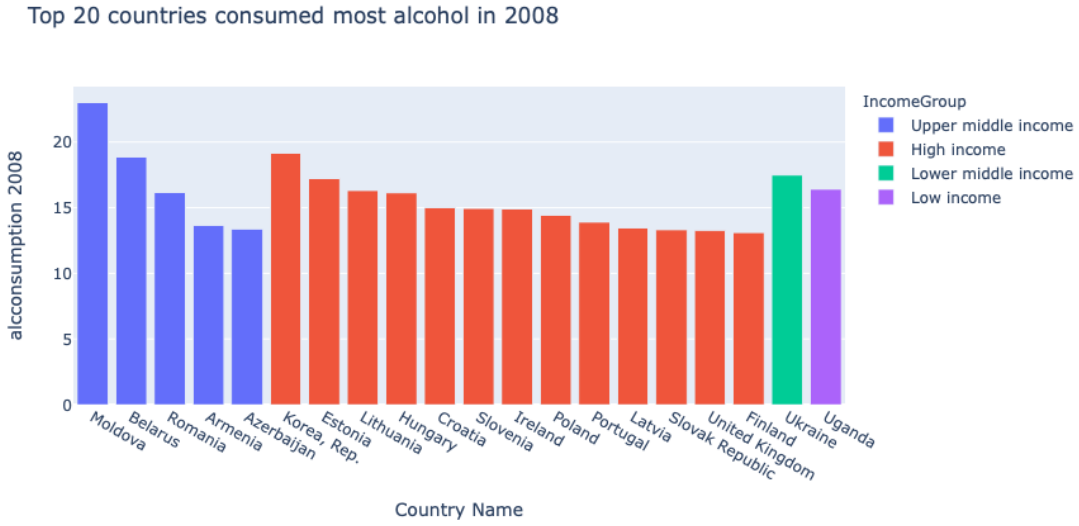

[15]:
```
# Wow, this is so much better!
# Finally, I will create a bubble map that show the total avg alcohol␣
 ↪consumption per capita. The only thing I could
# have add is to code 'year'as the animation frame but unfortunately our yearly␣
 ↪data is sporadic
combined_df['2008'] = combined_df['alcconsumption 2008'].fillna(0)
fig = px.scatter_geo(combined_df, locations='Country Code', color='IncomeGroup',
                     hover_name=combined_df.index, size='2008',
                    title='Bubble Graph of Global Alcohol Consumption by Income␣
 ↪Group in 2008')
fig.show()
```

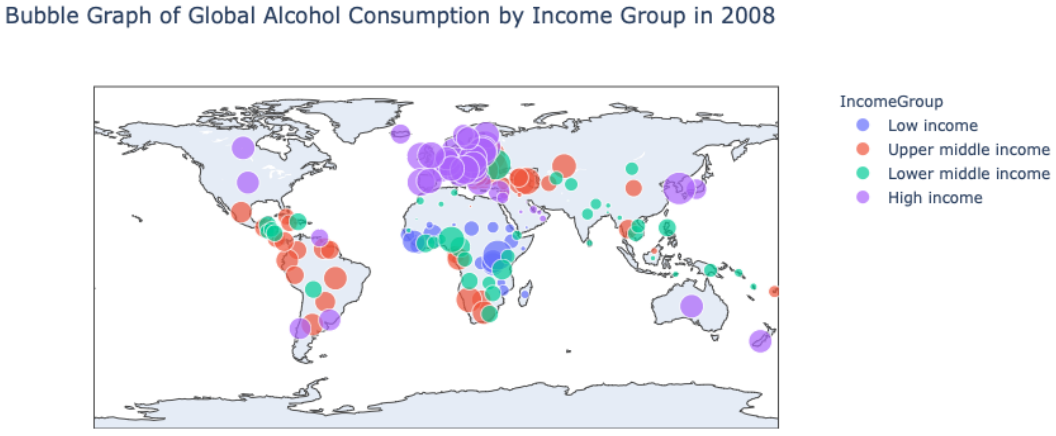

[16]:
```
# Bonus just for fun
# When I have more time to work on this, perhaps I want to look at something␣
 ↪more predictive than just simply
# summarizing data
# Moving on from descriptive statistics, I want to explore potential␣
 ↪relationships between these variables in the dataset
# Let's plot a few scatter plots for differnt pair of x and y axes

fig = px.scatter(combined_df, x='alcconsumption 2008', y='suicideper100th')
fig.show()
fig = px.scatter(combined_df, x='alcconsumption 2008', y='employrate')
fig.show()
```

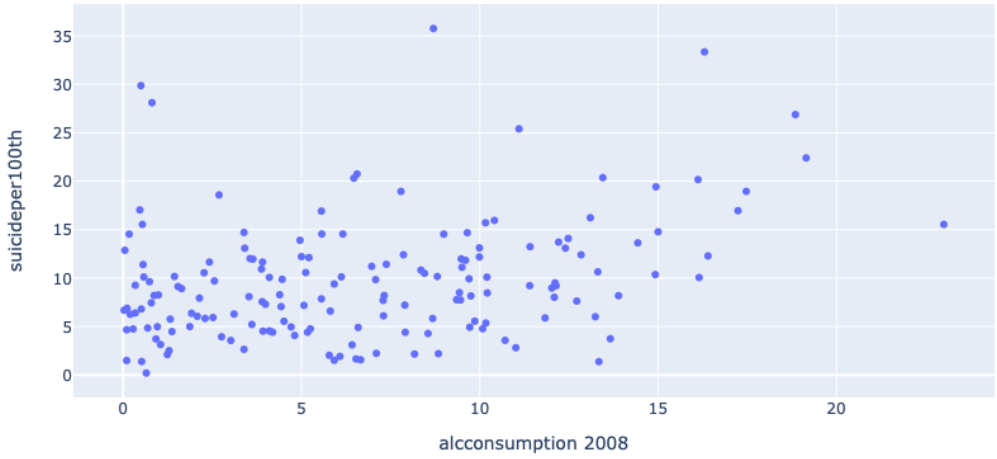

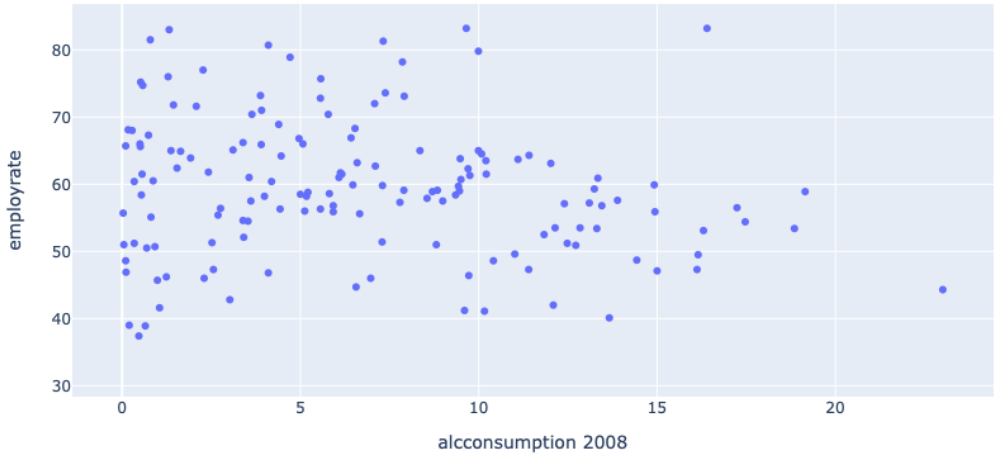

[17]:
```
# it looks like there might exist a positive correlation between alcohol␣
 ↪consumption and suicidal rate per 100,000th
# On the other hand, the relationship between alcohol consumption and␣
 ↪employment rate is not quiet clear
```